# Indexing

INF 551 Wensheng Wu

## Outline

- Types of indexes
- B+ trees

#### Indexes

- An <u>index</u> is a data structure that speeds up selections on the <u>search key field(s)</u>
- *Fields* = *attributes*
- Search key = any subset of the fields of a relation
  - Search key is not the same as key (minimal set of fields that uniquely identify a record in a relation).
- Entries in an index: (k, r), where:
  - k = key
  - r = record(s) OR record id(s)

### **Index Classification**

- Clustered/unclustered
  - Clustered = records sorted & stored in the order of search key
  - Unclustered = records are not sorted in key order
- B+ tree / hash table / ...

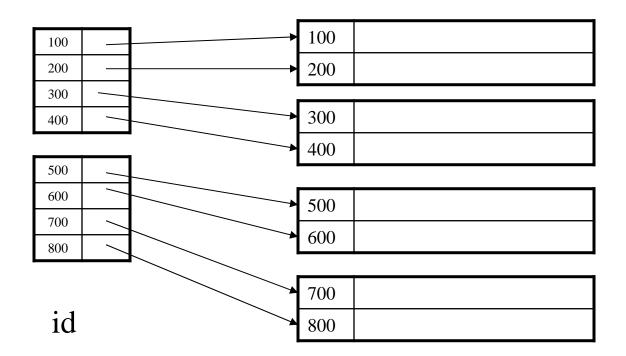
# MySQL

- Automatically creates a clustered index for:
  - Primary key if exists;
  - Otherwise 1<sup>st</sup> unique key;
  - If no unique keys, on row ID (a hidden attribute)
- Row data are stored with the clustered index

More details here

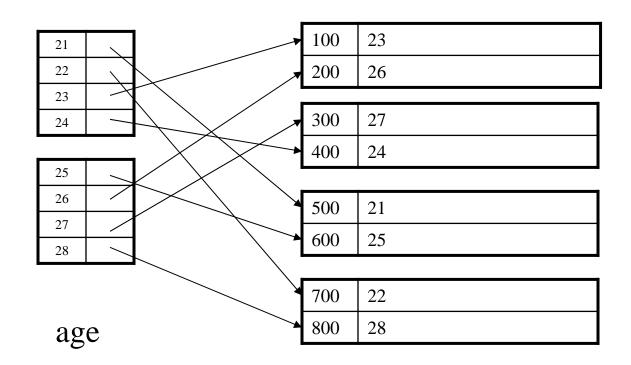
#### Clustered Index

- Records are sorted on the search key
  - E.g., employee(<u>id</u>, name, age, salary)



#### Unclustered Indexes

Records NOT sorted by the search key



Records sorted by id

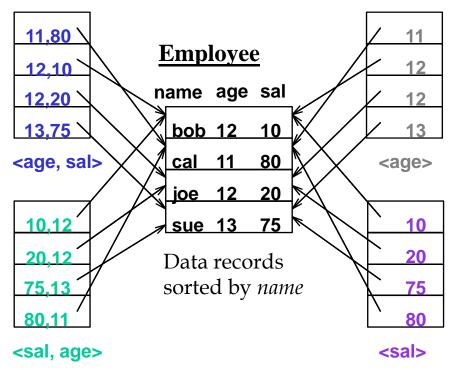
# Query Types

- Equality/point query: <attribute> = <value>
  - E.g., age = 20, sal = 75
- Range query: <attribute> <inequality operator> <value>
  - Inequality operator: <, >, <=, >=
  - E.g., age > 20 or sal <= 75

## Composite Search Keys

• *Composite Search Keys*: Search key = a list of fields.

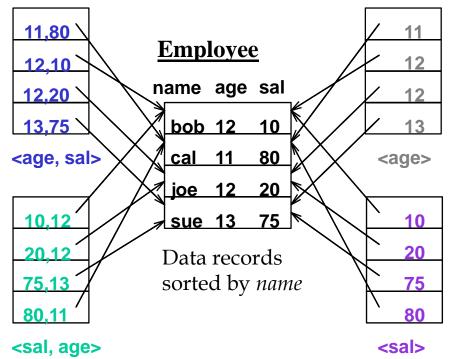
Keys in index sorted by <age,sal>: i.e., first by age; if ties, by sal



Keys sorted by <sal>

## Questions

- Which index is useful for queries:
  - -Sal > 75
  - Age = 12 and sal > 10
  - -Age > 12



## Outline

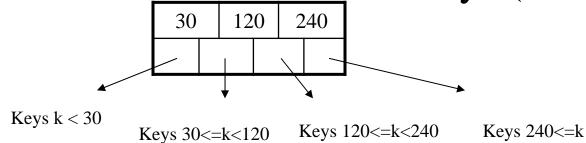
- Types of indexes
- B+ trees

#### B+ Trees

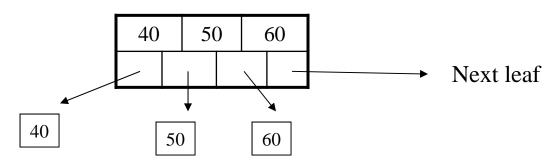
- Search trees
- Idea in B Trees:
  - make 1 node = 1 block
- Idea in B+ Trees:
  - Make leaves into a linked list
  - Efficiently support range queries

#### **B+ Trees Basics**

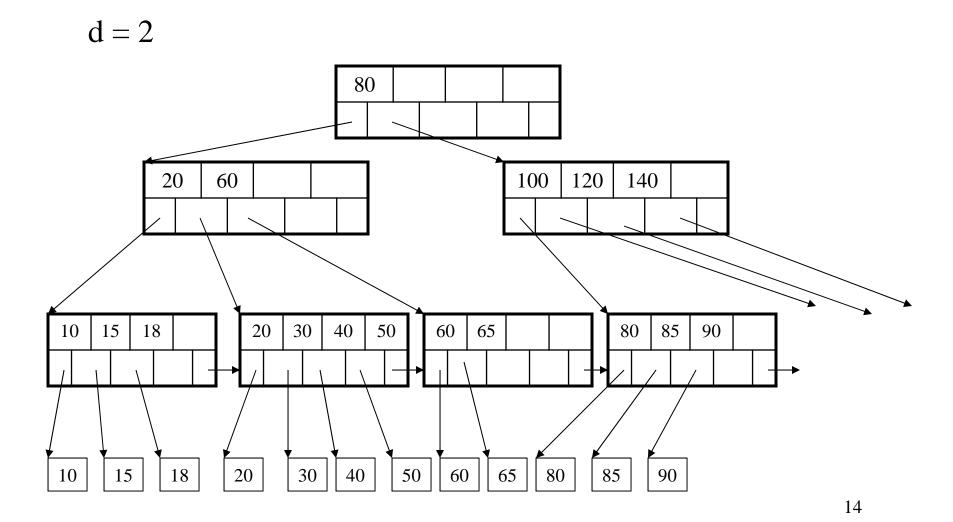
- Parameter d = the **degree** (also called order)
- Each node has >= d and <= 2d keys (except root)



• Each leaf has >=d and <= 2d keys:



# B+ Tree Example



## B+ Tree Design

- How large is d?
- Example:
  - Key size = 4 bytes
  - Pointer size = 8 bytes
  - Block size = 4096 byes
- 2d \* 4 + (2d+1) \* 8 <= 4096
- $d = 170 (\sim 170.33)$

#### B+ Trees in Practice

- Typical order d = 100.
- Typical fill-factor (minimum in practice): 66.7% (i.e., 2/3) (note minimum fill factor in design: 50%)
  - Minimum # of keys in a node = 133 (200 \* 2/3)
- Capacities (# of records which the index supports):
  - Height 1 (tree with a single root): 133 records
  - Height 2:  $133^2 = 17$ , 689 records (134\*133 to be exact)
  - Height 3:  $133^3 = 2$ , 352, 637 records ( $134^2 *133$ )
  - Height 4:  $133^4 = 312,900,721$  records  $(134^3 * 133)$

### B+-tree in Practice

• Can often hold top levels in buffer pool:

```
Level 1 = 1 page = 4KB
Level 2 = 133 pages = 532KB
Level 3 = 17,689 pages = 70, 756KB ~ 70MB
```

# Searching a B+ Tree

- Equality search:
  - Start at the root
  - Proceed down, to the leaf

Select name From people Where age = 25

- Range query [a, b]:
  - Finding the first leaf in the range
  - Then sequential traversal of leaves until ...

Select name
From people
Where 20 <= age
and age <= 30

# Searching a B+ Tree

- Range query [-, b]:
  - Finding the left-most leaf
  - Then sequential traversal of leaves until ...

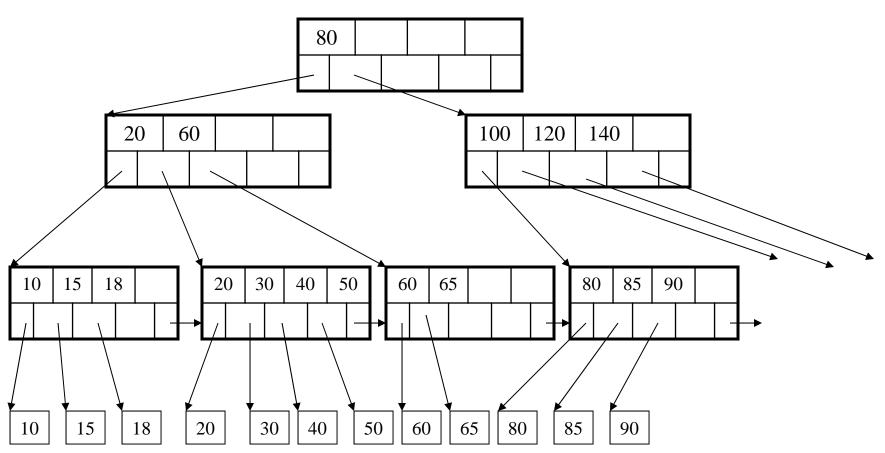
Select name From people Where age <= 30

- Range query [a, -]:
  - Finding the leaf with a
  - Then sequential traversal until ...

Select name From people Where 20 <= age

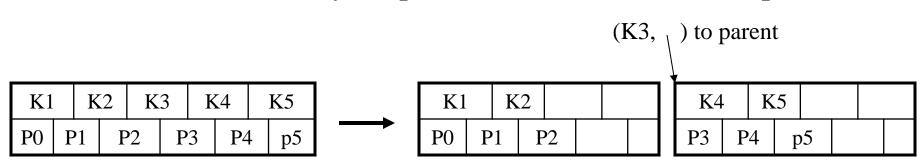
# Example

20 <= age and age <= 55



#### Insert (K, P)

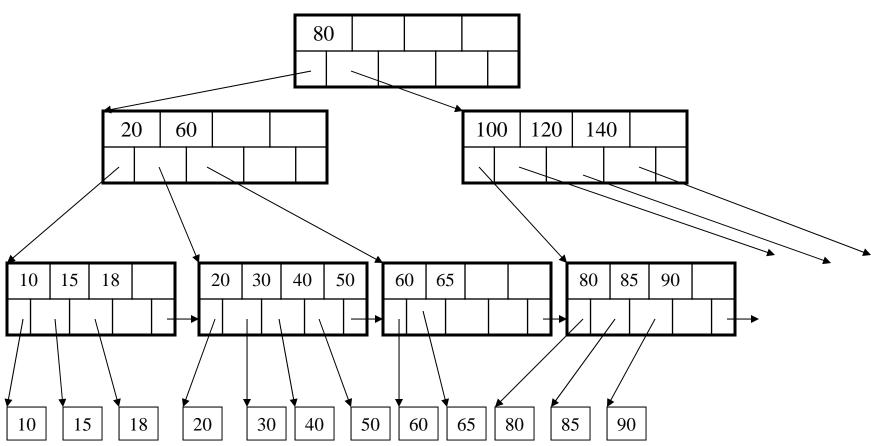
- Find leaf where K belongs, insert
- If no overflow (2d keys or less), stop
- If overflow (2d+1 keys), split node, insert middle into parent:



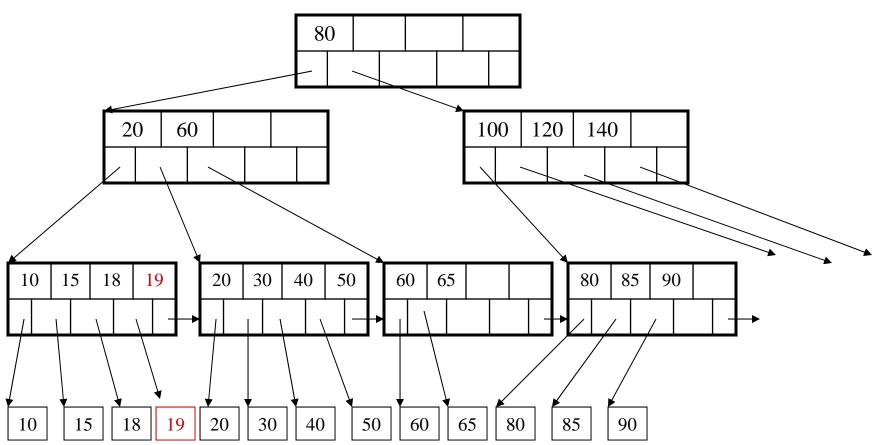
#### Notes

- Splitting of leaf may lead to splitting of its parent and ancestors
- When splitting a leaf, middle key (e.g., K3) is also kept in the new node on the right
- No need to retain middle key in the split node when splitting an internal node (but need to insert it into parent)
- When root is split, new root has only one key

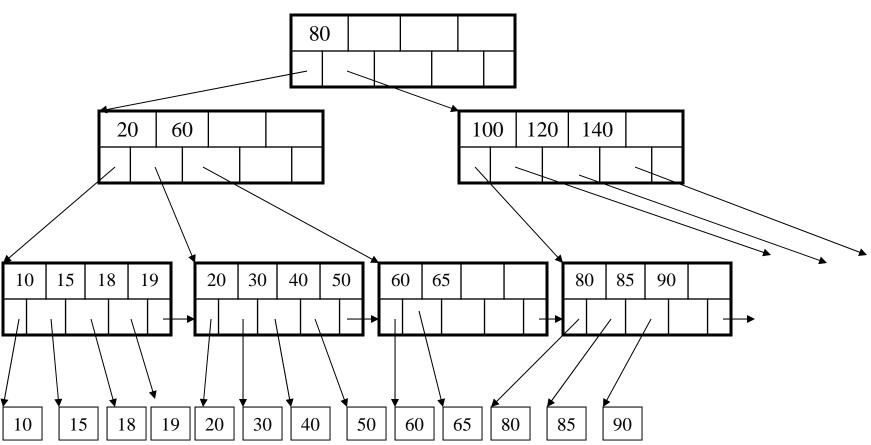




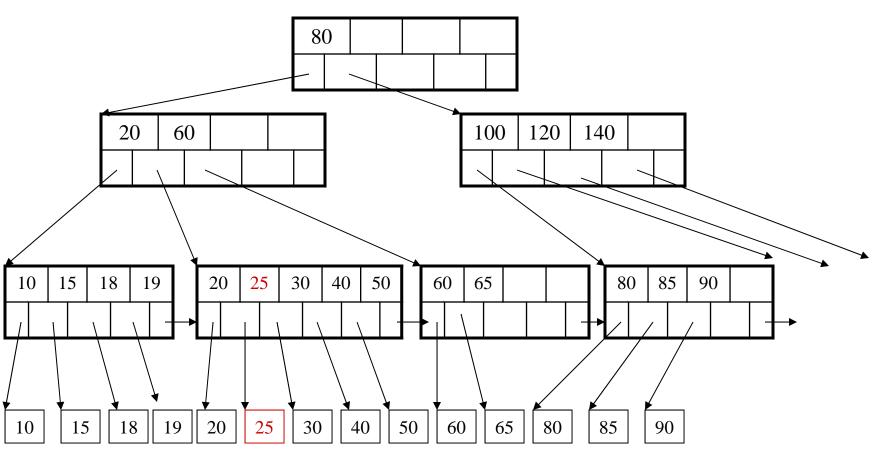
#### After insertion



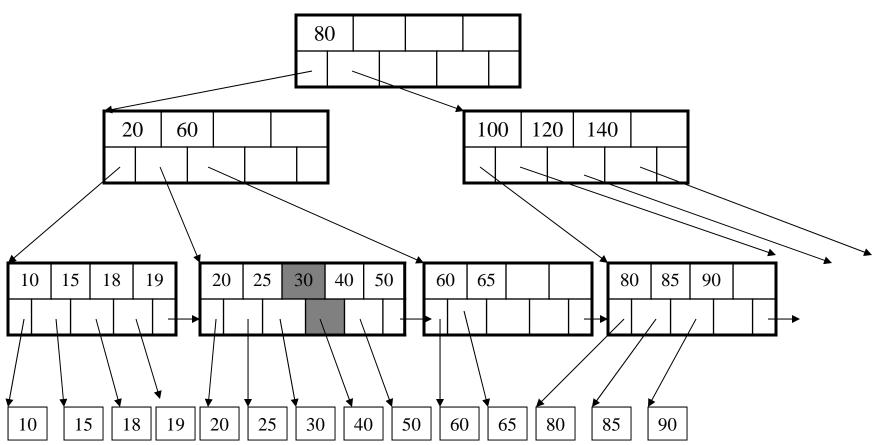
#### Now insert 25



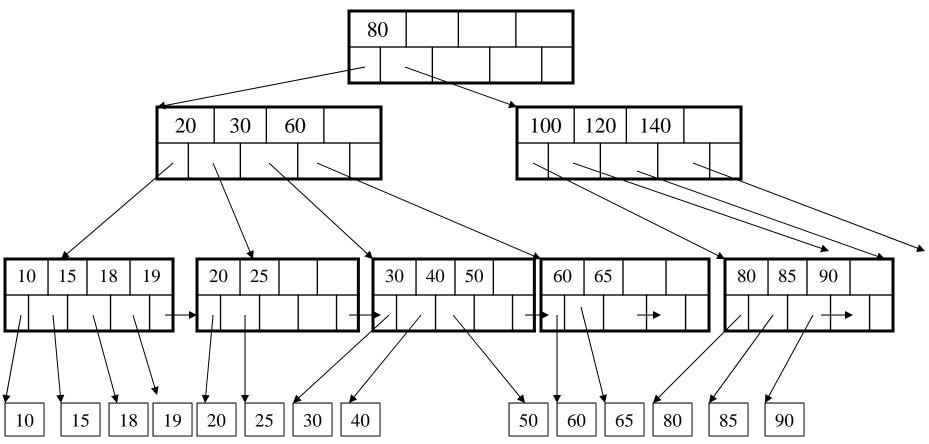
#### After insertion



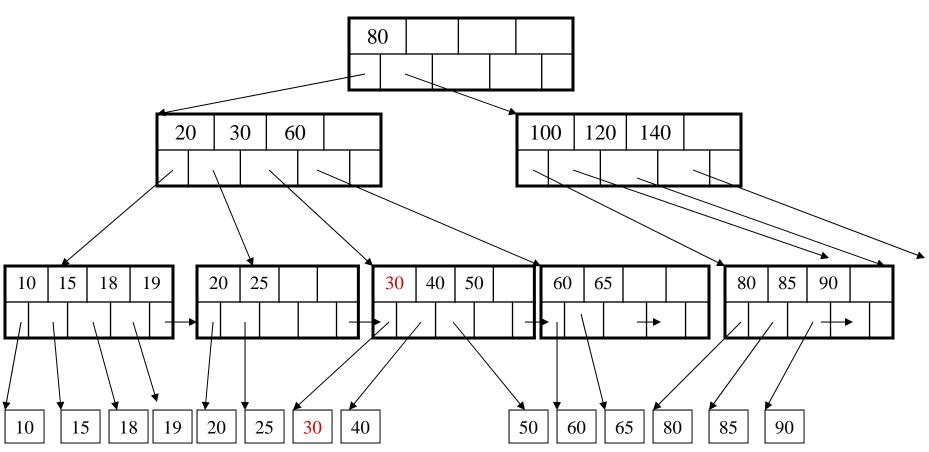
But now have to split!



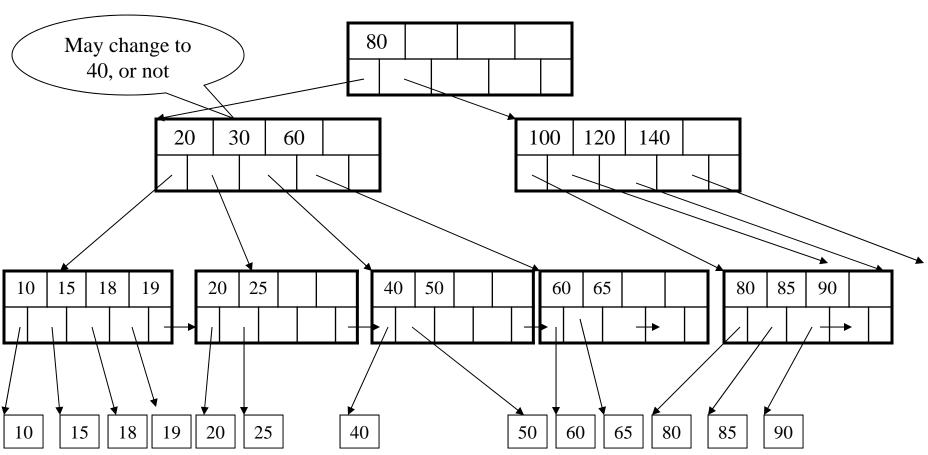
After the split



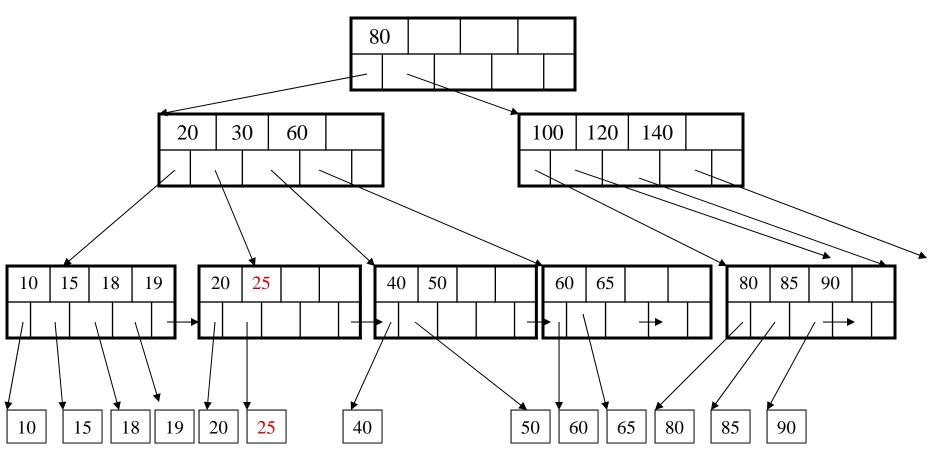
#### Delete 30

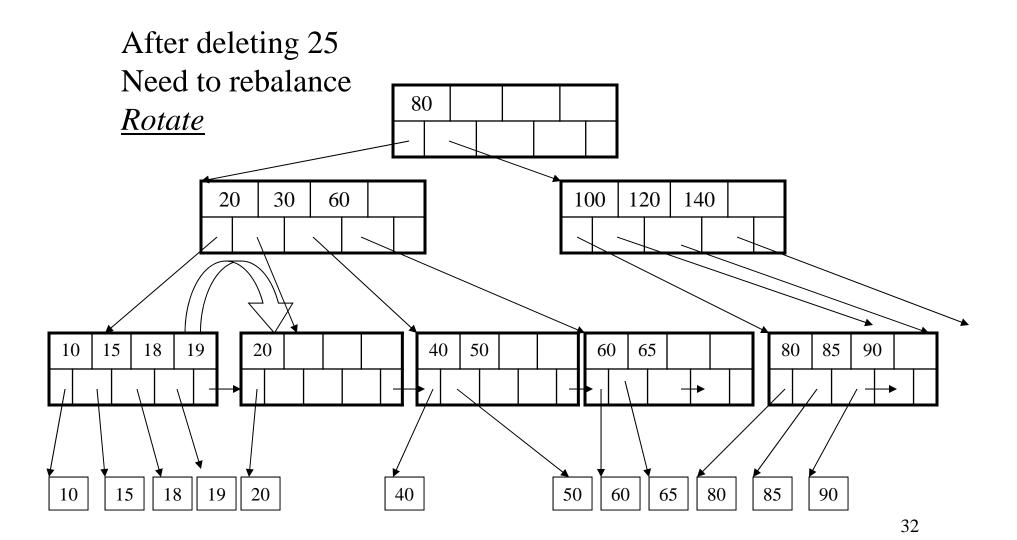




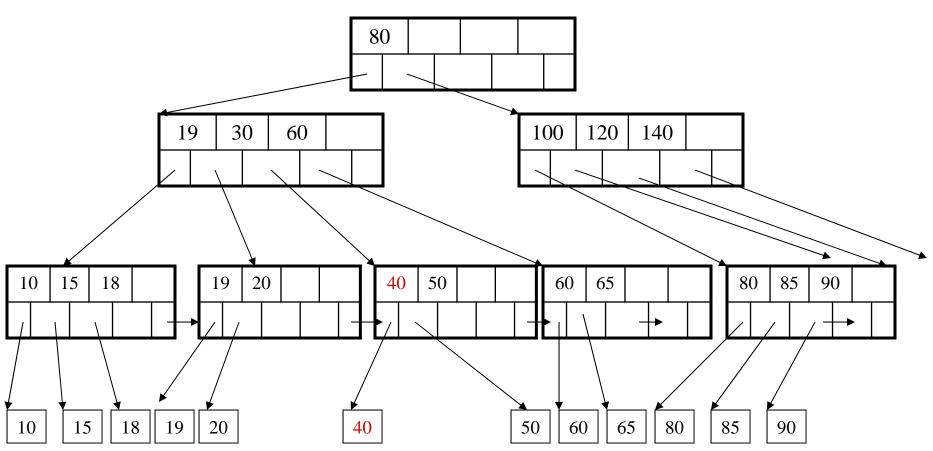


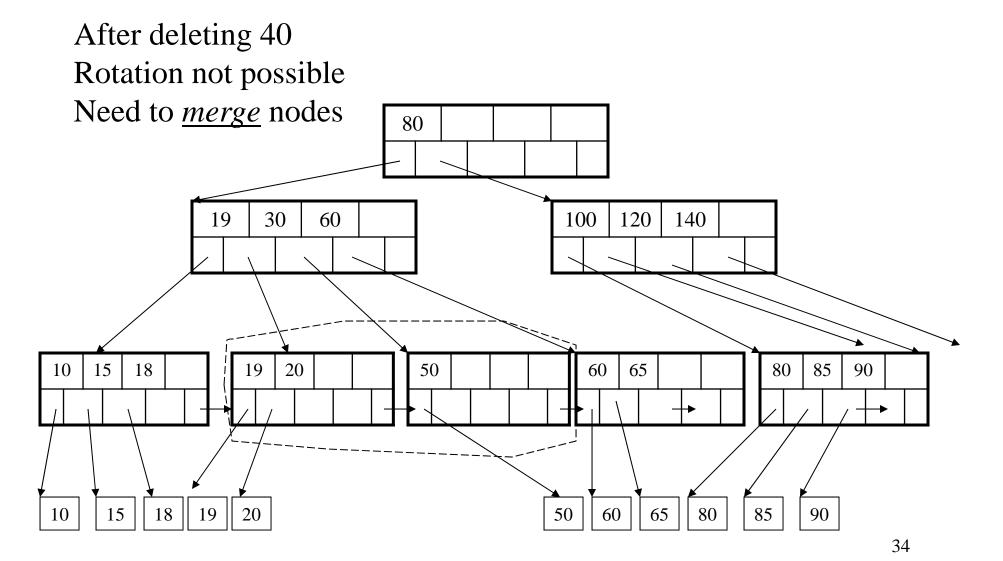
#### Now delete 25



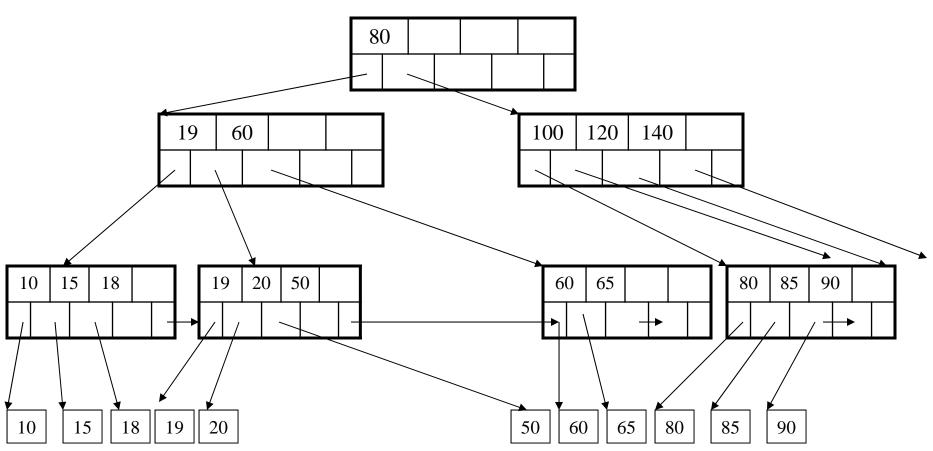


#### Now delete 40





#### Final tree



## **Deletion Strategy**

- If a node is below the min capacity after deletion...
- Try the following in the given order
  - 1. move a key from immediate left sibling;
  - 2. move a key from immediate right sibling;
  - 3. merge with immediate left sibling;
  - 4. merge with immediate right sibling
- Cases 3 and 4 may lead to further removal of key from parent, and more fixing

# Another insertion example

